Introduction to Machine Learning

Advanced Risk Minimization
Bernoulli Loss

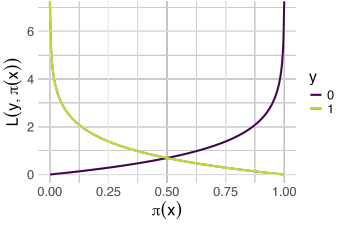

Learning goals

- Know the Bernoulli loss and related losses (log, logistic, binomial loss)
- Derive the risk minimizer
- Derive the optimal constant model

BERNOULLI LOSS ON PROBABILITIES

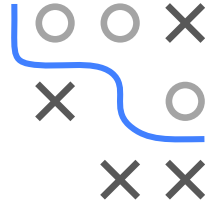

The likelihood function for n samples of a Bernoulli RV with probability π parametrized by $\boldsymbol{\theta}$ is

$$\mathcal{L}(\boldsymbol{\theta}) = \prod_{i=1}^{n} \pi\left(\mathbf{x}^{(i)} \mid \boldsymbol{\theta}\right)^{y^{(i)}} \left(1 - \pi\left(\mathbf{x}^{(i)} \mid \boldsymbol{\theta}\right)\right)^{1-y^{(i)}} \quad , \; y \in \{0, 1\}$$

Taking the negative log to turn products in into sums and maximization into minimization, we obtain the NLL:

$$-\ell(\boldsymbol{\theta}) = \sum_{i=1}^{n} -y^{(i)} \log\left(\pi\left(\mathbf{x}^{(i)} \mid \boldsymbol{\theta}\right)\right) - \left(1 - y^{(i)}\right) \log\left(1 - \pi\left(\mathbf{x}^{(i)} \mid \boldsymbol{\theta}\right)\right)$$

The Bernoulli loss is then defined as the NLL of a single sample

$$L\left(y, \pi(\mathbf{x})\right) = -y \log\left(\pi(\mathbf{x})\right) - (1 - y) \log\left(1 - \pi(\mathbf{x})\right) \quad , \; y \in \{0, 1\}$$

BERNOULLI LOSS LABEL ENCODING

For different label encoding $y' \in \{-1, +1\}$, recognize that $y = (y' + 1)/2$. Substituting y into $L(y, f(\mathbf{x}))$ yields

$$L(y, f(\mathbf{x})) = \log(1 + \exp(-y \cdot f(\mathbf{x}))) \quad \text{for } y \in \{-1, +1\}$$

- Two equivalent formulations for different label encodings
- All variants convex, differentiable

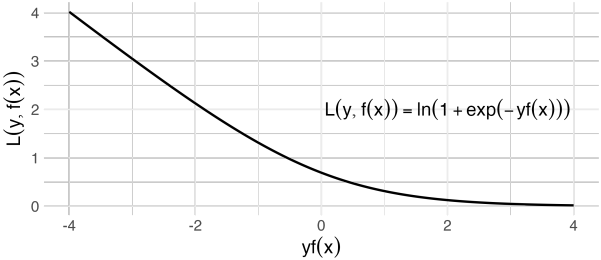

BERNOULLI LOSS: RISK MINIMIZER

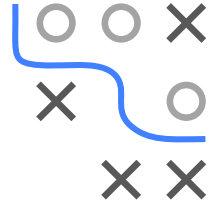

The risk minimizer for the Bernoulli loss defined for probabilistic classifiers $\pi(\mathbf{x})$ and on $y \in \{0, 1\}$ is

$$\pi^*(\tilde{\mathbf{x}}) = \eta(\tilde{\mathbf{x}}) = \mathbb{P}(y = 1 \mid \mathbf{x} = \tilde{\mathbf{x}}).$$

Proof: We can write the risk for binary y as follows:

$$\mathcal{R}(f) \quad = \quad \mathbb{E}_x \left[L(1, \pi(\mathbf{x})) \cdot \eta(\mathbf{x}) + L(0, \pi(\mathbf{x})) \cdot (1 - \eta(\mathbf{x})) \right],$$

with $\eta(\mathbf{x}) = \mathbb{P}(y = 1 \mid \mathbf{x})$ (see section on the 0-1-loss for more details).
For a fixed $\mathbf{x}$ we compute the point-wise optimal value c by setting the derivative to 0:

$$\begin{aligned}
\frac{\partial}{\partial c} \left(-\log c \cdot \eta(\mathbf{x}) - \log(1 - c) \cdot (1 - \eta(\mathbf{x})) \right) &= 0 \\
-\frac{\eta(\mathbf{x})}{c} + \frac{1 - \eta(\mathbf{x})}{1 - c} &= 0 \\
\frac{-\eta(\mathbf{x}) + \eta(\mathbf{x})c + c - \eta(\mathbf{x})c}{c(1 - c)} &= 0 \\
c &= \eta(\mathbf{x}).
\end{aligned}$$

BERNOULLI: OPTIMAL CONSTANT MODEL

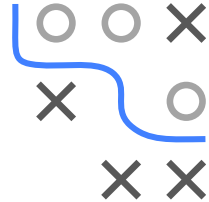

The optimal constant probability model $\pi(\mathbf{x}) = \theta$ w.r.t. the Bernoulli loss for labels from $\mathcal{Y} = \{0, 1\}$ is:

$$\hat{\theta} = \arg\min_{\theta} \mathcal{R}_{\text{emp}}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^{n} y^{(i)}$$

This is the fraction of class-1 observations in the observed data. We can prove this by setting the derivative of the risk to 0 and solving for θ. The optimal constant score model $f(\mathbf{x}) = \theta$ w.r.t. the Bernoulli loss labels from $\mathcal{Y} = \{-1, +1\}$ or $\mathcal{Y} = \{0, 1\}$ is:

$$\hat{\theta} = \arg\min_{\theta} \mathcal{R}_{\text{emp}}(\boldsymbol{\theta}) = \log \frac{n_+}{n_-} = \log \frac{n_+/n}{n_-/n}$$

where n_- and n_+ are the numbers of negative and positive observations. This shows a tight (and unsurprising) connection of this loss to log-odds.

Proving this is left as exercise.

BERNOULLI-LOSS: NAMING CONVENTION

We have seen several closely related loss functions:

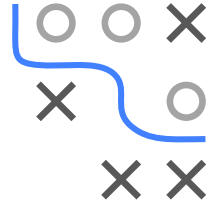

$$
\begin{aligned}
L(y, f(\mathbf{x})) &= \log(1 + \exp(-yf(\mathbf{x}))) \quad \text{for } y \in \{-1, +1\} \\
L(y, f(\mathbf{x})) &= -y \cdot f(\mathbf{x}) + \log(1 + \exp(f(\mathbf{x}))) \quad \text{for } y \in \{0, 1\} \\
L(y, \pi(\mathbf{x})) &= -y \log(\pi(\mathbf{x})) - (1 - y) \log(1 - \pi(\mathbf{x})) \quad \text{for } y \in \{0, 1\} \\
L(y, \pi(\mathbf{x})) &= -\frac{1 + y}{2} \log(\pi(\mathbf{x})) - \frac{1 - y}{2} \log(1 - \pi(\mathbf{x})) \quad \text{for } y \in \{-1, +1\}
\end{aligned}
$$

They are equally referred to as Bernoulli-, Binomial-, logistic-, log-, or cross-entropy loss.

OPTIMIZATION PROPERTIES: CONVERGENCE

In case of **complete separation**, optimization might fail, e.g.:

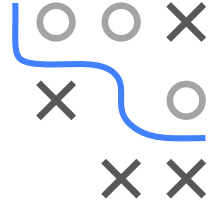

- **Bernoulli loss** is margin-based and strictly monotonically decreasing in $y \cdot f(\mathbf{x})$:

$$L\left(y, f(\mathbf{x})\right) = \log\left(1 + \exp\left(-yf(\mathbf{x})\right)\right)$$

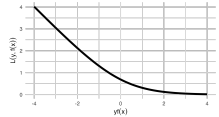

- f linear in $\boldsymbol{\theta}$, e.g., **logistic regression** with $f(\mathbf{x} \mid \boldsymbol{\theta}) = \boldsymbol{\theta}^\top \mathbf{x}$
- Data perfectly separable by our learner, so we can find $\boldsymbol{\theta}$:

$$y^{(i)} f\left(\mathbf{x}^{(i)} \mid \boldsymbol{\theta}\right) = y^{(i)} \boldsymbol{\theta}^T \mathbf{x}^{(i)} > 0 \;\; \forall \mathbf{x}^{(i)}$$

- Can now a construct a strictly better $\boldsymbol{\theta}$

$$\mathcal{R}_{\text{emp}}(2 \cdot \boldsymbol{\theta}) = \sum_{i=1}^{n} L\left(2y^{(i)} \boldsymbol{\theta}^T \mathbf{x}^{(i)}\right) < \mathcal{R}_{\text{emp}}(\boldsymbol{\theta})$$

- As $||\boldsymbol{\theta}||$ increases, sum strictly decreases, as argument of L is strictly larger
- We can iterate that, so there is no local (or global) optimum, and no numerical procedure can converge