


Introduction to Machine Learning

Information Theory

Entropy I





Learning goals

- Entropy measures expected information for discrete RVs
- Know entropy and its properties

INFORMATION THEORY


- **Information Theory** is a field of study based on probability theory.
- Foundation was laid by Claude Shannon in 1948; since then been applied in: communication theory, computer science, optimization, cryptography, machine learning and statistical inference.
- Quantify the "amount" of information gained or uncertainty reduced when a random variable is observed.
- Also about storing and transmitting information.



ENTROPY AS SURPRISAL AND UNCERTAINTY

For a discrete random variable X with domain $\mathcal{X} \ni x$ and pmf $p(x)$:

$$\begin{aligned} H(X) &:= H(p) = -\mathbb{E}[\log_2(p(X))] &= -\sum_{x \in \mathcal{X}} p(x) \log_2 p(x) \\ &= \mathbb{E} \left[\log_2 \left(\frac{1}{p(X)} \right) \right] &= \sum_{x \in \mathcal{X}} p(x) \log_2 \frac{1}{p(x)} \end{aligned}$$



Some technicalities first:



- H is actually Greek capital letter **E**ta (η) for **e**ntropy
- Base of the log simply specifies the unit we measure information in, usually bits (base 2) or 'nats' (base e)
- If $p(x) = 0$ for an x , then $p(x) \log_2 p(x)$ is taken to be zero, because $\lim_{p \rightarrow 0} p \log_2 p = 0$.

ENTROPY AS SURPRISAL AND UNCERTAINTY

$$H(X) = -\mathbb{E}[\log_2(p(X))] = -\sum_{x \in \mathcal{X}} p(x) \log_2 p(x)$$

Now: What's the point?

- The negative log probabilities $-\log_2 p(x)$ are called "surprisal"
- More surprising means less likely
- PMFs surprising, so with higher H , when events more equally likely
- Entropy is simply expected surprisal




- The final entropy is $H(X) = 2.12$ (bits).

ENTROPY BASIC PROPERTIES

$$H(X) := H(p) = -\mathbb{E}[\log_2(p(X))] = -\sum_{x \in \mathcal{X}} p(x) \log_2 p(x)$$

- ❶ Entropy is non-negative, so $H(X) \geq 0$
- ❷ If one event has probability $p(x) = 1$, then $H(X) = 0$
- ❸ Adding or removing an event with $p(x) = 0$ doesn't change it
- ❹ $H(X)$ is continuous in probabilities $p(x)$


All these properties follow directly from the definition.




ENTROPY RE-ORDERING

- 5 Symmetry. If the values $p(x)$ in the pmf are re-ordered, entropy does not change (proof is trivial).


Entropy $H(p) = 2.12$





Entropy $H(p) = 2.12$



Entropy $H(p) = 2.12$





Entropy $H(p) = 2.12$



ENTROPY OF UNIFORM DISTRIBUTIONS






Let X be a uniform, discrete RV with g outcomes (g -sided fair die).

$$H(X) = - \sum_{i=1}^g \frac{1}{g} \log_2 \left(\frac{1}{g} \right) = \log_2 g$$



The more sides a die has, the harder it is to predict the outcome.
Unpredictability grows *monotonically* with the number of potential outcomes, but at a decreasing rate.

ENTROPY IS MAXIMAL FOR UNIFORM



- Naive observation:
Entropy min for 1-point and max for uniform distribution

ENTROPY IS MAXIMAL FOR UNIFORM

- ⑥ Entropy is maximal for a uniform distribution,
for domain of size g : $H(X) \leq -g \frac{1}{g} \log_2(\frac{1}{g}) = \log_2(g)$.

Proof: So we want to maximize w.r.t. all p_i :

$$\operatorname{argmax}_{p_1, p_2, \dots, p_g} - \sum_{i=1}^g p_i \log_2 p_i$$

subject to

$$\sum_{i=1}^g p_i = 1$$

