



Introduction to Machine Learning

Boosting

Gradient Boosting: CWB Basics 1



Learning goals

- Concept of CWB
- Which base learners do we use
- Built-in feature selection

COMPONENTWISE GRADIENT BOOSTING I


GB (with trees), has strong predictive performance but is difficult to interpret unless the base learners are stumps.

The aim of CWB is to find a model that exhibits:

- strong predictive performance,
- interpretable components,
- automatic selection of components,
- is sparser than a model fitted with maximum-likelihood estimation.

This is achieved by using “nice” base learners which yield familiar statistical models in the end.



Because of this, CWB is also often referred to as **model-based boosting**.



BASE LEARNERS I

In GB only one kind of base learner $\mathcal{B}$ is used, e.g., regression trees.

For CWB we generalize this to multiple base learner sets $\{\mathcal{B}_1, \dots, \mathcal{B}_J\}$ with associated parameter spaces $\{\Theta_1, \dots, \Theta_J\}$, where $j \in \{1, 2, \dots, J\}$ indexes the type of base learner.




In each iteration, base learners are fitted to the **pseudo residuals** $\tilde{r}^{[m]}$.


BASE LEARNERS III

Two BLs of the same type can simply be added by adding up their parameter vectors:

$$b_j(\mathbf{x}, \theta^{[1]}) + b_j(\mathbf{x}, \theta^{[2]}) = b_j(\mathbf{x}, \theta^{[1]} + \theta^{[2]}).$$



Thus, if $\{b_j(\mathbf{x}, \theta^{[1]}), b_j(\mathbf{x}, \theta^{[2]})\} \in \mathcal{B}_j$, then $b_j(\mathbf{x}, \theta^{[1]} + \theta^{[2]}) \in \mathcal{B}_j$.




COMPONENTWISE BOOSTING ALGORITHM I

Different from GB, multiple base learners $b_j \in \mathcal{B}_j, j = 1, \dots, J$, are fitted and only best-fitting one is selected and updated.

Algorithm Componentwise Gradient Boosting.

- 1: Initialize $f^{[0]}(\mathbf{x}) = \arg \min_{\theta_0 \in \mathbb{R}} \sum_{i=1}^n L(y^{(i)}, \theta_0)$
 - 2: **for** $m = 1 \rightarrow M$ **do**
 - 3: For all i : $\tilde{r}^{[m](i)} = - \left[\frac{\partial L(y, f)}{\partial f} \right]_{f=f^{[m-1]}(\mathbf{x}^{(i)})}$
 - 4: **for** $j = 1 \rightarrow J$ **do**
 - 5: Fit regression base learner $b_j \in \mathcal{B}_j$ to the vector of pseudo-residuals $\tilde{r}^{[m]}$:
 - 6: $\hat{\theta}_j^{[m]} = \arg \min_{\theta \in \Theta_j} \sum_{i=1}^n (\tilde{r}^{[m](i)} - b_j(\mathbf{x}^{(i)}, \theta))^2$
 - 7: **end for**
 - 8: $j^{[m]} = \arg \min_j \sum_{i=1}^n (\tilde{r}^{[m](i)} - \hat{b}_j(\mathbf{x}^{(i)}, \hat{\theta}_j^{[m]}))^2$
 - 9: Update $f^{[m]}(\mathbf{x}) = f^{[m-1]}(\mathbf{x}) + \alpha \hat{b}_{j^{[m]}}(\mathbf{x}, \hat{\theta}_{j^{[m]}}^{[m]})$
 - 10: **end for**
 - 11: Output $\hat{f}(\mathbf{x}) = f^{[M]}(\mathbf{x})$
-

(Same as for GB, New inner loop for CWB)



COMPONENTWISE BOOSTING ALGORITHM I

Algorithm Componentwise Gradient Boosting (inner loop).


4: **for** $j = 1 \rightarrow J$ **do**

5: Fit regression base learner $b_j \in \mathcal{B}_j$ to the vector of pseudo-residuals $\tilde{r}^{[m]}$:




6: $\hat{\theta}_j^{[m]} = \arg \min_{\theta \in \Theta_j} \sum_{i=1}^n (\tilde{r}^{[m](i)} - b_j(\mathbf{x}^{(i)}, \theta))^2$

7: **end for**

8: $j^{[m]} = \arg \min_j \sum_{i=1}^n (\tilde{r}^{[m](i)} - \hat{b}_j(\mathbf{x}^{(i)}, \hat{\theta}_j^{[m]}))^2$



Iteration $m, j = 1, \sum_{i=1}^n (\tilde{r}^{[m](i)} - \hat{b}_1(x_1^{(i)}, \hat{\theta}_1^{[m]}))^2 = 24.4$:



COMPONENTWISE BOOSTING ALGORITHM I

Algorithm Componentwise Gradient Boosting (inner loop).


4: **for** $j = 1 \rightarrow J$ **do**

5: Fit regression base learner $b_j \in \mathcal{B}_j$ to the vector of pseudo-residuals $\tilde{r}^{[m]}$:




6: $\hat{\theta}_j^{[m]} = \arg \min_{\theta \in \Theta_j} \sum_{i=1}^n (\tilde{r}^{[m](i)} - b_j(\mathbf{x}^{(i)}, \theta))^2$

7: **end for**

8: $j^{[m]} = \arg \min_j \sum_{i=1}^n (\tilde{r}^{[m](i)} - \hat{b}_j(\mathbf{x}^{(i)}, \hat{\theta}_j^{[m]}))^2$



Iteration $m, j = 2, \sum_{i=1}^n (\tilde{r}^{[m](i)} - \hat{b}_2(x_2^{(i)}, \hat{\theta}_2^{[m]}))^2 = 43.2$:



COMPONENTWISE BOOSTING ALGORITHM I

Algorithm Componentwise Gradient Boosting (inner loop).


4: **for** $j = 1 \rightarrow J$ **do**

5: Fit regression base learner $b_j \in \mathcal{B}_j$ to the vector of pseudo-residuals $\tilde{r}^{[m]}$:




6: $\hat{\theta}_j^{[m]} = \arg \min_{\theta \in \Theta_j} \sum_{i=1}^n (\tilde{r}^{[m](i)} - b_j(\mathbf{x}^{(i)}, \theta))^2$

7: **end for**

8: $j^{[m]} = \arg \min_j \sum_{i=1}^n (\tilde{r}^{[m](i)} - \hat{b}_j(\mathbf{x}^{(i)}, \hat{\theta}_j^{[m]}))^2$



Iteration $m, j = 3, \sum_{i=1}^n (\tilde{r}^{[m](i)} - \hat{b}_3(x_3^{(i)}, \hat{\theta}_3^{[m]}))^2 = 35.2$:



A 3x3 grid with a blue path starting at the top-left cell (row 1, column 1) and ending at the bottom-right cell (row 3, column 3). The path moves right to (1,2), down to (2,2), and then down to (3,3). The cells (1,3), (2,1), and (3,1) are marked with 'X', while the other cells are empty.

$$8: j^{[m]} = \arg \min_j \sum_{i=1}^n (\tilde{r}^{[m](i)} - \hat{b}_j(\mathbf{x}^{(i)}, \hat{\theta}_j^{[m]}))^2$$

FEATURE SELECTION IN CWB I

In CWB, we often define BLs on a single feature

$$b_j(x_j, \theta) \quad \text{for } j = 1, 2, \dots, p.$$

Allows natural form of feature selection:

- When we select the best BL in one iter of CWB, we thereby also only select one (associated) feature
- Note that a feature (or rather a BL associated with it) can be selected in multiple iters, so $\leq M$ features are selected

